




Designing an Azure Data Solution Crash Course

Microsoft Certified: Azure Data Engineer Associate

December/2020



Reza Salehi

Cloud Consultant and Trainer


O'Reilly Media Instructor, Pluralsight Author



@zaalion



2008 - 2018



Course Overview

Questions & Resources

- You can post questions in the QnA panel
- Resources are in the course repository
 - <https://github.com/zaalion/oreilly-dp-200-201>
- Reach out:
 - Twitter: [@zaalion](https://twitter.com/zaalion)



DP-201 Skills Measured

Exam DP-201: Designing an Azure Data Solution skills



DP-201 Skills Measured

- Skills measured:
 - Design Azure data storage solutions (40-45%)
 - Design data processing solutions (25-30%)
 - Design for data security and compliance (25-30%)



DP-201 Candidate Profile

- Microsoft Azure data engineers
 - Collaborate with business stakeholders to identify and meet the data requirements.
 - To design data solutions that use Azure data services.
 - Design vs. implement



Azure Data Engineers

- Responsible for data-related design tasks
- Include designing Azure data storage solutions
- Use relational and non-relational data stores, batch and real-time data processing solutions
- Data security and compliance solutions



DP-201 Candidates

- Must design data solutions that use:
 - Azure Cosmos DB
 - Azure Synapse Analytics
 - Azure Data Lake Storage, Azure Data Factory, Azure Stream Analytics, Azure Databricks, and Azure Blob storage.



Design Azure Data Storage Solutions

Design Azure Data Storage Solutions

- Recommend an Azure data storage solution based on requirements
- Design non-relational cloud data stores
- Design relational cloud data stores





Big Data

Extremely large data sets which can be analyzed to find patterns and trends.





Data Lake

- A data store
- The data is raw
- The data purpose is completely clear yet
- Heterogeneous data is stored in one place (structured, unstructured, text, Access, Excel, binary, json, AVRO, etc.)
- Data is not strongly typed





Data Warehouse

- Used for reporting and data analysis
- Data is structured
- Data is generally pre-processed (e.g., using Databricks, HDInsight, ML, etc.)
- Optimized for reporting (e.g., not too many joins)
- The data purpose is clear



Recommend an Azure Data Storage Solution Based on Requirements

- Choose the correct data storage solution to meet the technical and business requirements
- Choose the partition distribution type



Choosing the Right Data Storage

- Relational databases
- Document databases
- Key/Value databases
- Graph databases
- Column family databases
- Object storage
- File share
- Data analytics databases
- Search Engine databases
- Time Series databases



Choosing the Right Data Storage

- Store logs / Azure services' output
 - Azure Blob Storage
- Low latency document database
 - Azure Cosmos DB Core API
- Database for social media
 - Azure Cosmos DB Graph API
- Migrating from MongoDB
 - Azure Cosmos MongoDB API



Choosing the Right Data Storage

- Building search around your existing data
 - Azure Cognitive Search
- Fast cache store
 - Azure Cache for Redis (Azure Redis)
- Highly relational data
 - Azure SQL Database
 - Other relational options
- Cheap column database
 - Azure Table Storage





Choosing the Right Data Storage

- Structured data
 - Azure SQL Database, MySQL, PostgreSQL, MariaDB
- Unstructured data
 - Azure Cosmos DB, Azure Table Storage
- Blobs / files
 - Azure Blob Storage, Data Lake Gen 2



Why Partition Your Data?

- Data partitioning
 - Improve scalability
 - Improve performance
 - Improve security
 - Provide operational flexibility
 - Match the data store to the pattern of use
 - Improve availability



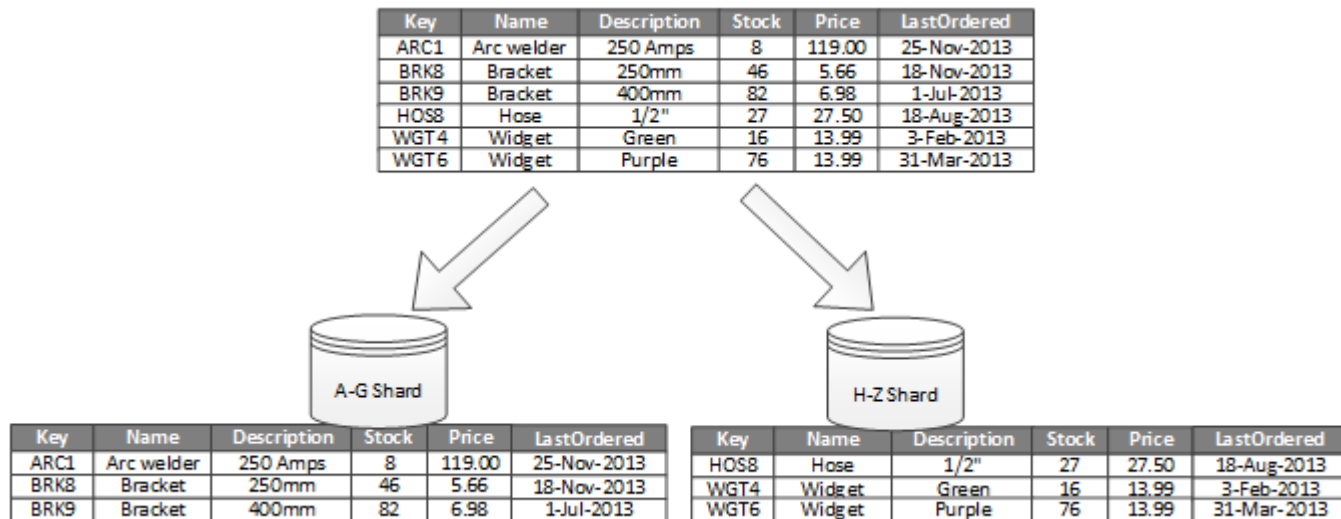


Choose the Partition Distribution Type

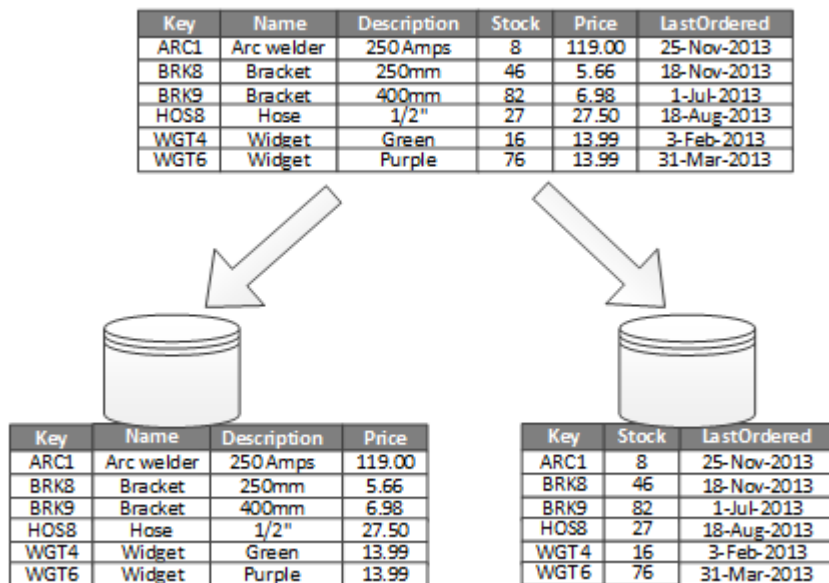
- Data partitioning types
 - Horizontal
 - Vertical
 - Functional



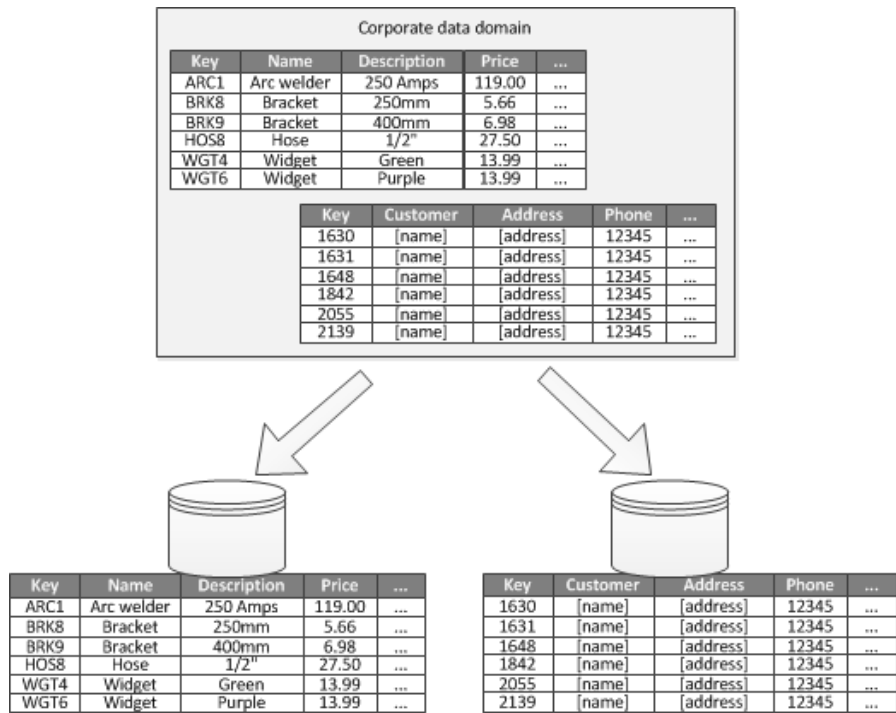
Horizontal Partitioning (Sharding)



Vertical Partitioning



Functional Partitioning



<https://docs.microsoft.com/en-us/azure/architecture/best-practices/data-partitioning#functional-partitioning>






Data Partitioning

- Azure Cosmos DB
- Azure Table Storage
- Azure Blob Storage
- Azure SQL Database
- Other services




Demo

- Azure Blob Storage
- Azure Data Lake Gen 2
- Azure Table Storage
- Azure File Storage
- Azure Queue Storage



Demo

- Azure SQL Database
- Other relational options (IaaS and PaaS)



Design Non-relational Cloud Data Stores

- Design a solution that uses Cosmos DB, Data Lake Storage Gen2, or Blob storage
- Select the appropriate Cosmos DB API
- Design data distribution and partitions
- Design for scale (including multi-region, latency, and throughput)
- Design a disaster recovery strategy
- Design for high availability





NoSQL Databases

- Data is stored by means other than related tables
- Document databases
- Key-value databases
- Wide-column databases
- Graph databases





Select the Appropriate Cosmos DB API

- Cosmos DB APIs
 - Azure Cosmos DB SQL API
 - Azure Cosmos DB's API for MongoDB
 - Azure Cosmos DB Cassandra API
 - Azure Cosmos DB Gremlin API
 - Azure Cosmos DB Table API
 - vs. Azure Table Storage





Cosmos DB Data Distribution

- Cosmos DB Data Distribution
 - Azure Cosmos DB multi-homing APIs
 - Consistency levels in Azure Cosmos DB

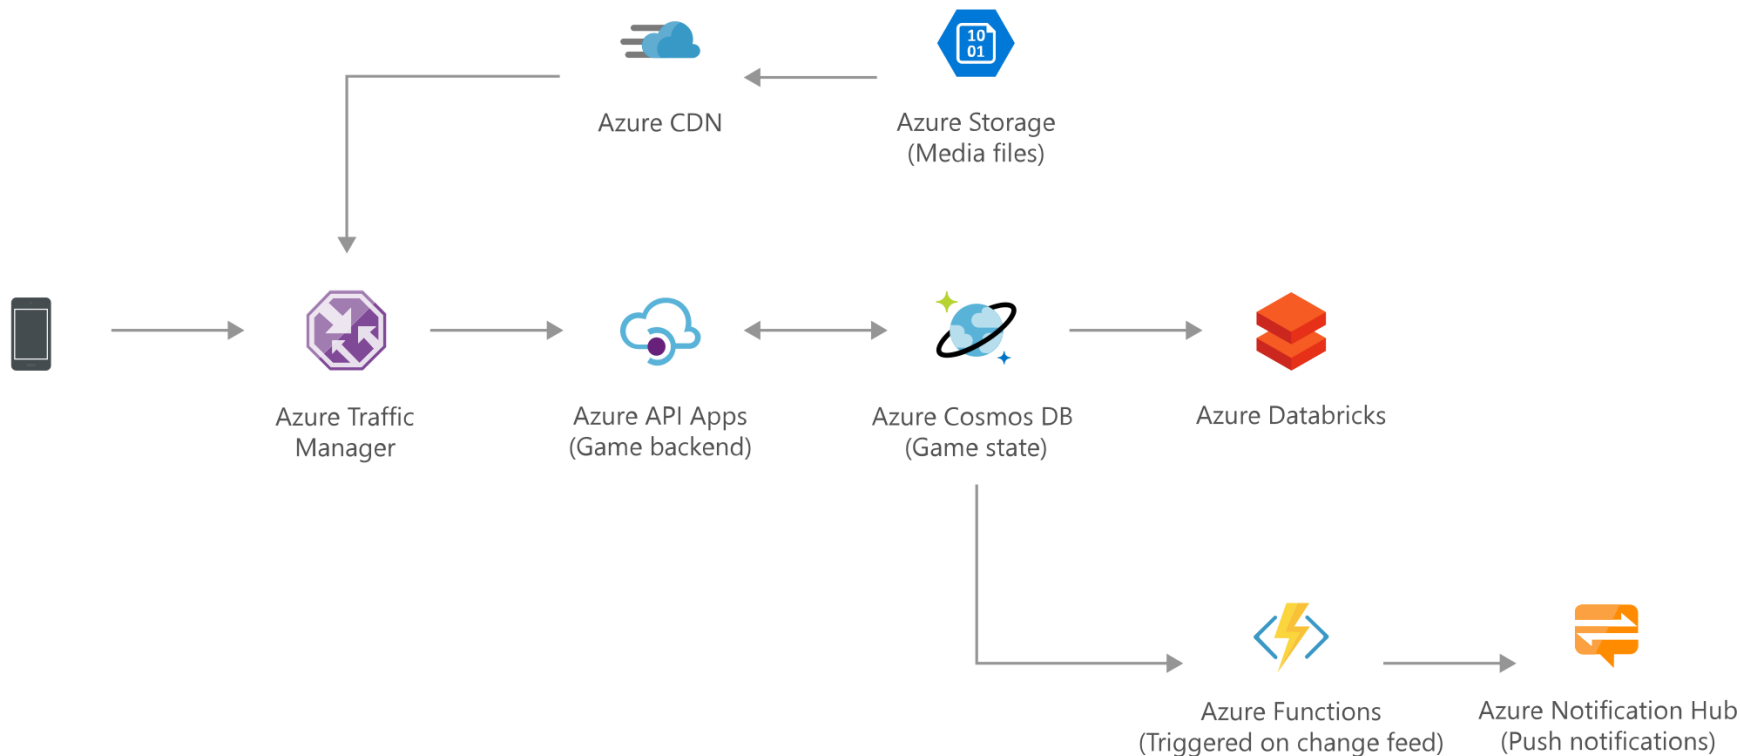


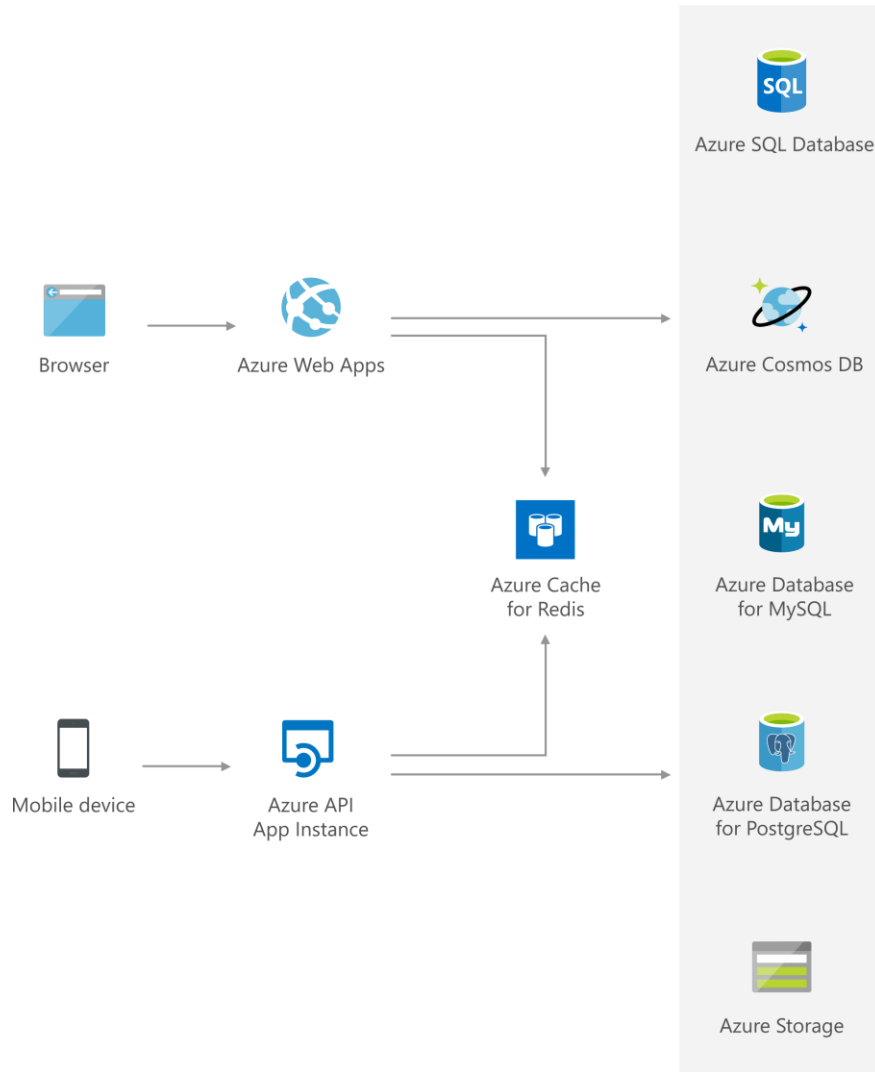
Azure Architectures

- <https://docs.microsoft.com/en-us/azure/architecture/browse/#databases>



Azure Storage and Cosmos DB





Data Cache



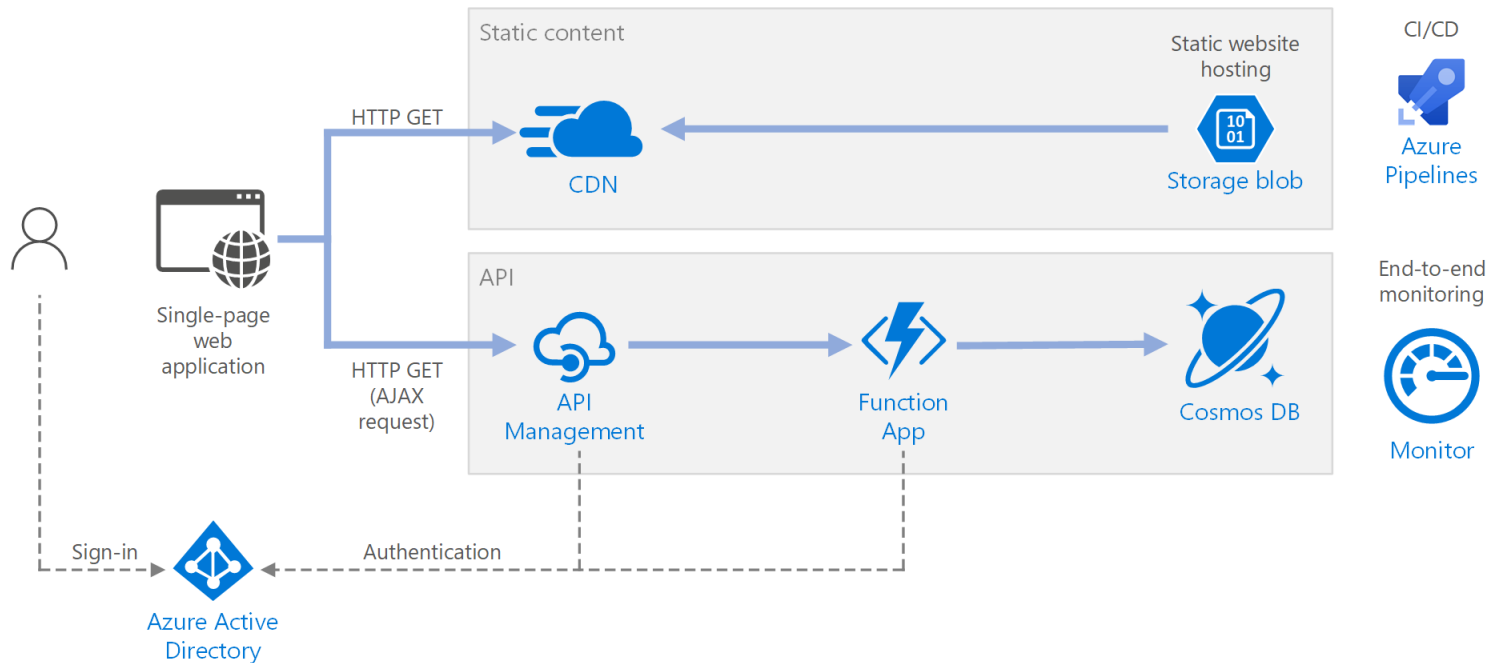


Design a Solution That Uses Data Lake Storage Gen2 & Blobs

- <https://docs.microsoft.com/en-us/azure/architecture/browse/#storage>



Azure Storage





Azure Data Lake Storage Gen 2

- Is built on top of Azure blob storage, simply enable a flag
- Hadoop compatible file system, HDFS (hierarchical namespace, folders)
- Store structured and non-structured data
- The purpose of the stored data is yet to be determined
- Great features of Azure Storage such as lifecycle management and hot/cold tiers
- Used as a data source for services such as Azure Databricks and Azure Synapse




High Availability

- Azure Storage Account
- Azure Cosmos DB




Demo

- Azure Cosmos DB
 - Provisioning
 - Multiple APIs
 - Data distribution
 - Multi-region
 - Adjusting throughput




Demo

- Azure Cosmos DB and other Azure services



Demo



- Azure Storage Account and other Azure services



Design Relational Cloud Data Stores

- Design data distribution and partitions
- Design for scale (including latency, and throughput)
- Design a solution that uses Azure Synapse Analytics
- Design a disaster recovery strategy
- Design for high availability





Azure Data Warehouse

- Ability to pause and resume (to save cost)
- Massive parallel processing
- Massive horizontal scaling (vs. vertical in Azure SQL Database)
- OLAP (vs. OLTP in Azure SQL Database)
- PolyBase T-SQL queries



PolyBase

- PolyBase enables SQL Server (and Synapse) to process Transact-SQL queries that read data from external data sources.
- Such as CSV files, etc.



Azure Synapse Analytics

- Components:
 - Synapse SQL: Complete T-SQL based analytics
 - Dedicated SQL pool (pay per DWU provisioned)
 - Serverless SQL pool (pay per TB processed)
 - Spark: Deeply integrated Apache Spark
 - Synapse Pipelines: Hybrid data integration (e.g., Azure Data Factory)
 - Studio: Unified user experience



Azure Synapse Analytics

Dedicated SQL pool (formerly SQL DW)



Dedicated SQL pool

Azure Synapse Analytics



Dedicated SQL pools



Serverless SQL pool



Apache Spark pools

Pipelines (Data Integration)

Shared metadata system

Common security model

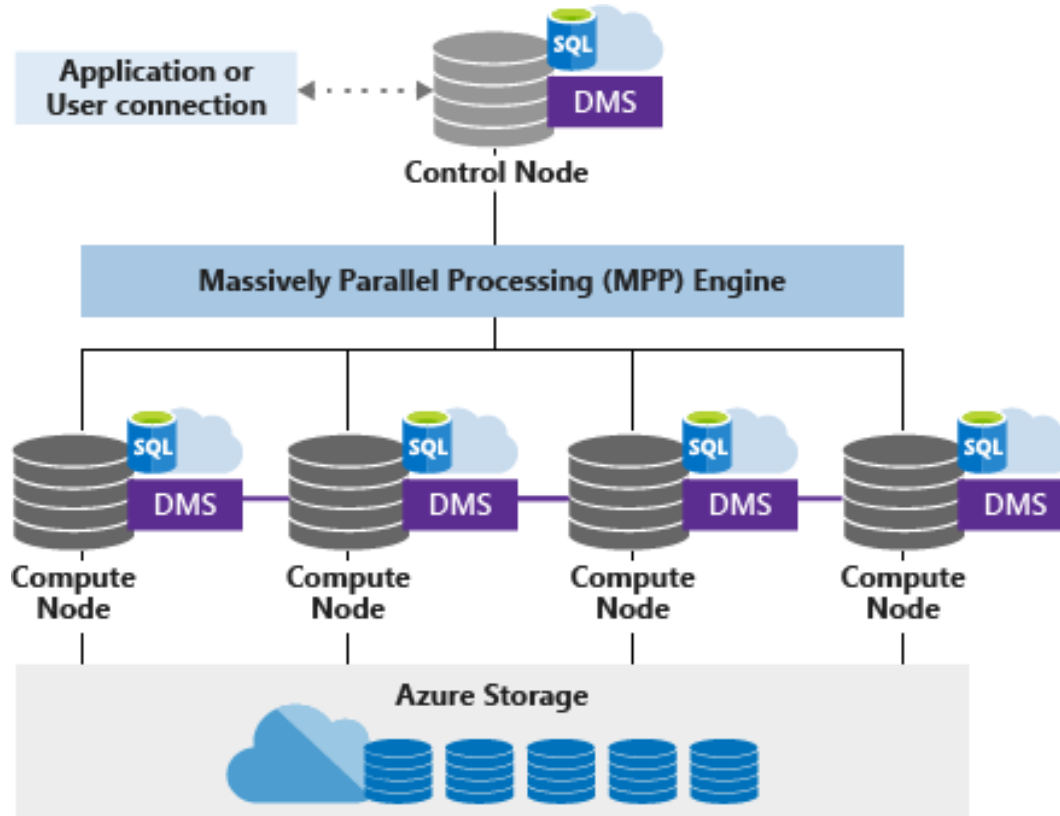
Connected Services

Synapse workspace

Synapse Studio



Azure Synapse Analytics






Azure Synapse Analytics Storage


- Azure Synapse Analytics Storage sharding options:
 - Hash-distributed tables
 - Round-robin distributed tables
 - Replicated Tables



Synapse Hash-distributed Tables



Synapse Replicated Tables



Synapse Round-robin Distributed Tables

- Delivers fast performance when used as a staging table for loads
- Distributes data evenly across the table but without any further optimization.
- It is quick to load data into a round-robin table
- Query performance can often be better with hash distributed tables.
- Joins on round-robin tables require reshuffling data, which takes additional time.



Azure Synapse Analytics Storage

- CREATE TABLE (Azure Synapse Analytics)




Design a Solution That Uses Azure Synapse

- <https://docs.microsoft.com/en-us/azure/architecture/browse/#databases>
- <https://docs.microsoft.com/en-us/azure/architecture/browse/#integration>




Azure Synapse Analytics



<https://docs.microsoft.com/en-us/azure/architecture/reference-architectures/data/enterprise-bi-synapse>



Azure Synapse Analytics



<https://docs.microsoft.com/en-us/azure/architecture/reference-architectures/data/enterprise-bi-adf>



Azure Synapse Studio

- [Open Synapse Studio](#)



All services > Azure Synapse Analytics > syn-demo01

Azure Synapse Ana... <<

zaalion (Default Directory)

+ Add ⚙️ Manage view ▾ ⋮

Filter by name...

Name ↑↓

syn-demo01 ⋮



syn-demo01 | Firewalls

Synapse workspace

🔍 Search (Ctrl+/)



💾 Save

✕ Discard

+ Add client IP

Settings

👤 SQL Active Directory admin

📊 Properties

🔒 Locks

Analytics pools

🗄️ SQL pools

⚙️ Apache Spark pools

Security

🛡️ Encryption

🔒 Firewalls

👤 Managed identities

🔗 Private endpoint connections

🔑 Approved Azure AD tenants

📄 Azure SQL Auditing

📘 The IPs listed below will have full access to Synapse workspace 'syn-demo01'.

Allow Azure services and resources to access this workspace

ON

OFF

Client IP address

99.230.107.165

Rule name

Start IP

End IP

allowAll

0.0.0.0


255.255.255.255 ⋮

High Availability

- Azure SQL Database
- Azure Synapse



Demo



- Azure Synapse



Design Data Processing Solutions

Design Data Processing Solutions

- Design batch processing solutions
- Design real-time processing solutions



Design Batch Processing Solutions

- Design batch processing solutions that use Data Factory and Azure Databricks
- Identify the optimal data ingestion method for a batch processing solution
- Identify where processing should take place, such as at the source, at the destination,
or in transit



Batch Processing vs. Stream Processing

- Batch processing: working with stored data (e.g. logs, historical data)
- Stream processing: working with live incoming data (e.g. sensors, IoT devices, audit logs)






Batch Processing

The process of loading data from a source, processing it and loading the result into a target data store. (ETL & ELT)









Batch Processing



Stream Processing

Ingest

-  IoT Devices
-  Logs, Files
-  Customer data, Financial transactions
-  Weather data
-  Business Apps

-  Event Hubs
-  Azure blob storage
-  IoT Hub

Analyze

Continuous Intelligence/Real-time analytics







Stream Analytics

Reference Data
SQL DB, Blob store

Real-time scoring
Azure ML service

Deliver

-  Alerts and actions
Event Hubs, Service Bus, Azure Functions etc
-  Dynamic Dashboarding
Power BI
-  Data Warehousing
Azure Synapse Analytics
-  Storage/ Archival
SQL DB, Azure Data Lake Gen 1 & Gen 2, Cosmos DB, Blob storage, etc




Databricks

- Cloud-hosted Apache Spark platform.
- Microsoft partnered with Databricks and introduced Azure Databricks
- Azure Databricks is a data analytics platform optimized for the Microsoft Azure cloud services platform.



Azure Databricks

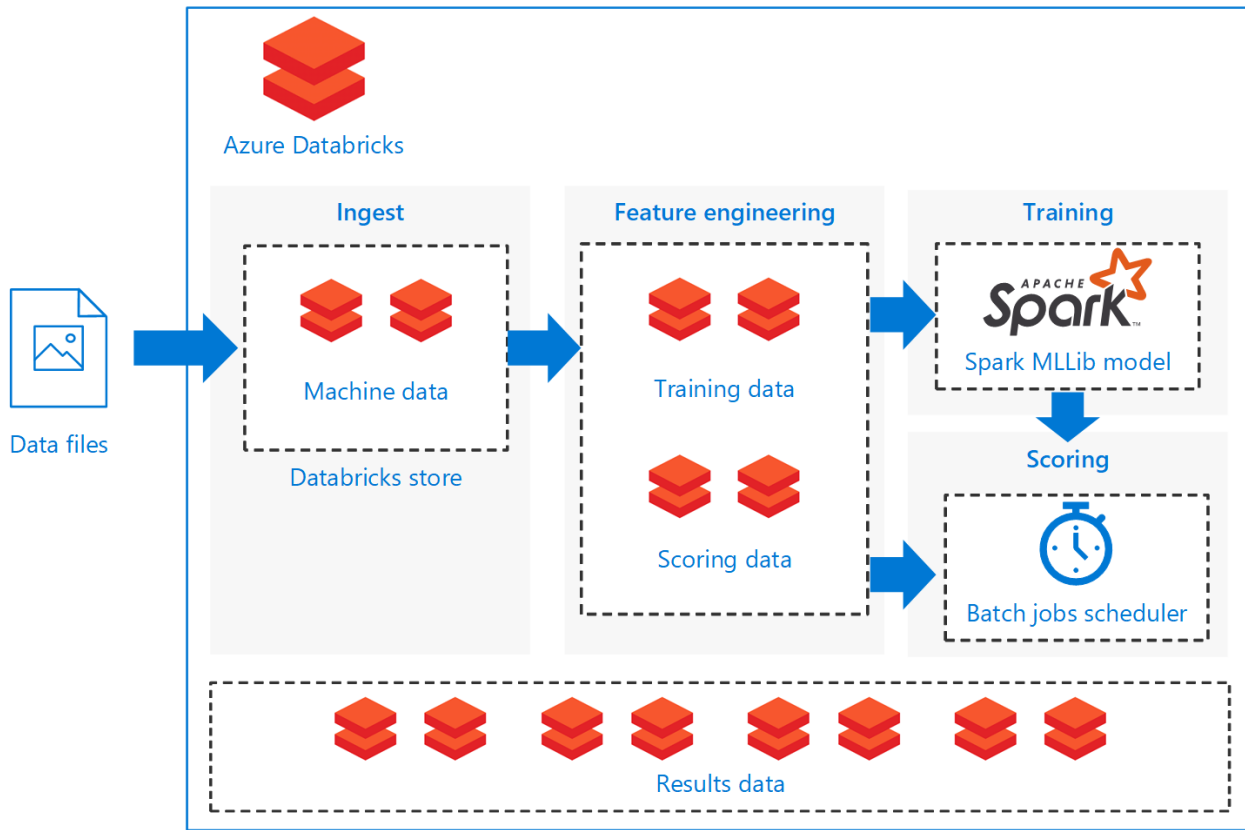


Design a Solution That Uses Azure Databricks


- <https://docs.microsoft.com/en-us/azure/architecture/browse/#ai--machine-learning>




Azure Databricks



Azure Databricks



Azure Databricks



All services > Azure Synapse Analytics > syn-demo01


Azure Synapse Ana... <<

zaalion (Default Directory)

+ Add ⚙ Manage view ▾ ⋮

Filter by name...

Name ↑↓

 syn-demo01 ⋮



syn-demo01 | Firewalls

Synapse workspace

🔍 Search (Ctrl+/)





💾 Save


✕ Discard

+ Add client IP


Settings


 SQL Active Directory admin

 Properties


 Locks

Analytics pools


 SQL pools


 Apache Spark pools


Security


 Encryption


 Firewalls

 Managed identities

 Private endpoint connections

 Approved Azure AD tenants

 Azure SQL Auditing

 The IPs listed below will have full access to Synapse workspace 'syn-demo01'.

Allow Azure services and resources to access this workspace

ON

OFF

Client IP address

99.230.107.165

Rule name

Start IP

End IP

allowAll

0.0.0.0

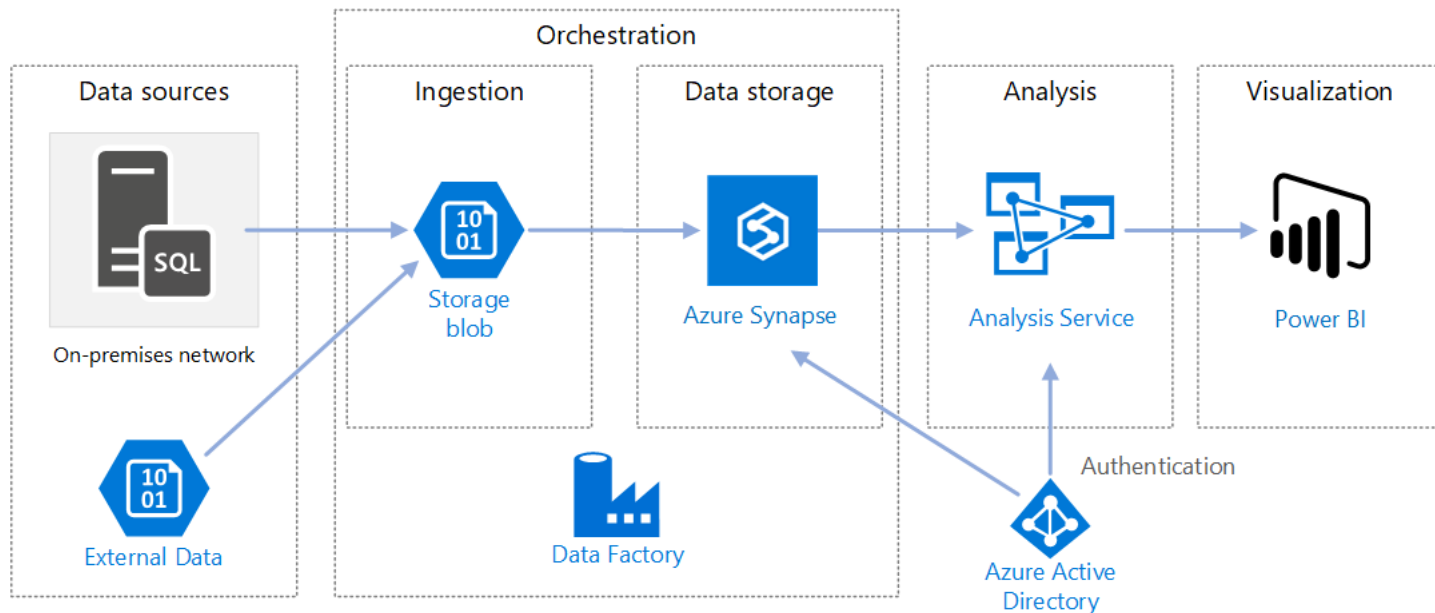
255.255.255.255 ⋮

Design a Solution That Uses Azure Data Factory

- <https://docs.microsoft.com/en-us/azure/architecture/browse/#analytics>



Azure Data Factory




<https://docs.microsoft.com/en-us/azure/architecture/reference-architectures/data/enterprise-bi-adf>



Demo

- Azure Data Factory

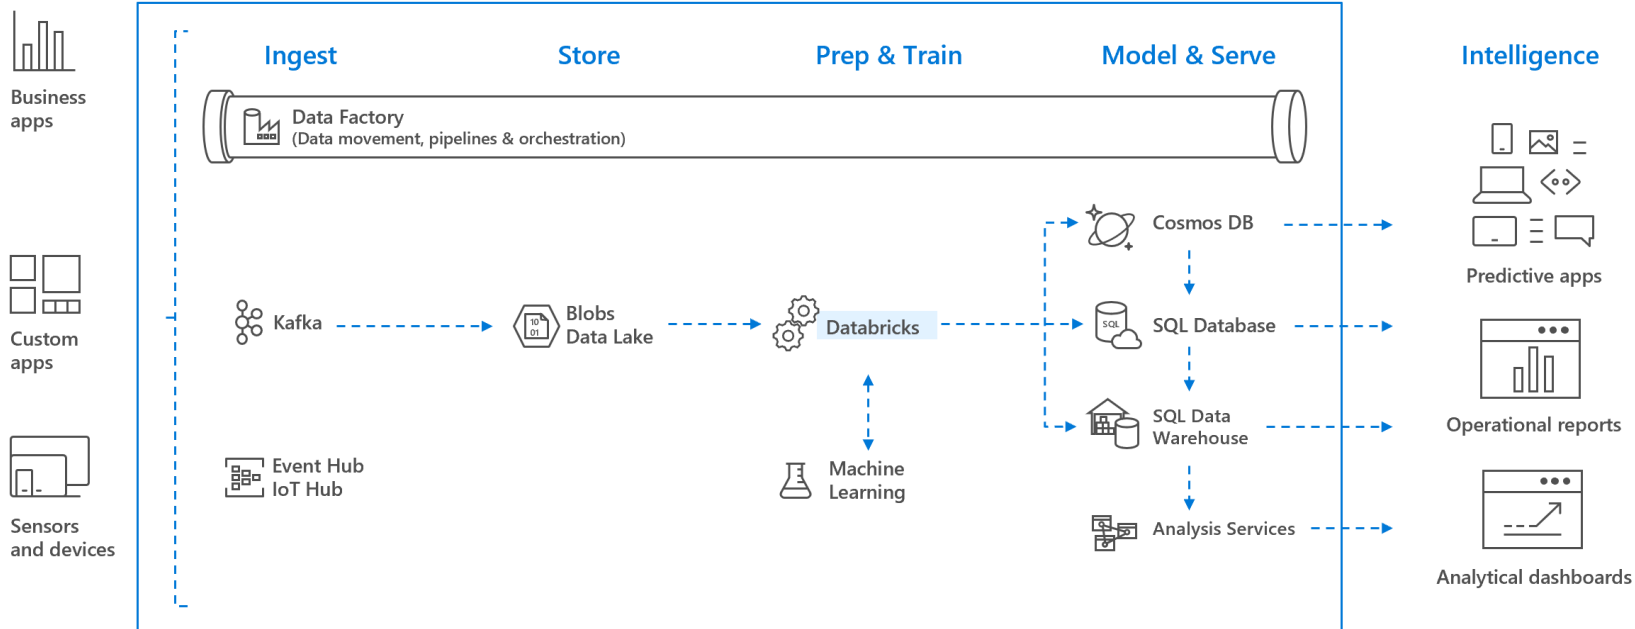


Design real-time processing solutions

- Design for real-time processing by using Stream Analytics and Azure Databricks
- Design and provision compute resources



Azure Databricks Workspace




<https://docs.microsoft.com/en-us/azure/databricks/scenarios/what-is-azure-databricks-ws>








Demo

- Azure Databricks



Azure Stream Analytics

Ingest

-  IoT Devices
-  Logs, Files
-  Customer data, Financial transactions
-  Weather data
-  Business Apps



Event Hubs



Azure blob storage



IoT Hub

Analyze

Continuous Intelligence/Real-time analytics



Stream Analytics



Reference Data
SQL DB, Blob store



Real-time scoring
Azure ML service

Deliver



Alerts and actions

Event Hubs, Service Bus,
Azure Functions etc



Dynamic Dashboarding

Power BI



Data Warehousing

Azure Synapse
Analytics



Storage/ Archival

SQL DB, Azure Data Lake Gen 1 &
Gen 2, Cosmos DB, Blob storage, etc



Develop Streaming Solutions

- Azure Stream Analytics
 - Ingest and process real-time data
 - Ingest from IoT Hub, Event Hubs and *Blob Storage*
 - Process using a *SQL-like* language
 - Output to several services such as *Event Hubs*, *Power BI*, Logic Apps, etc.

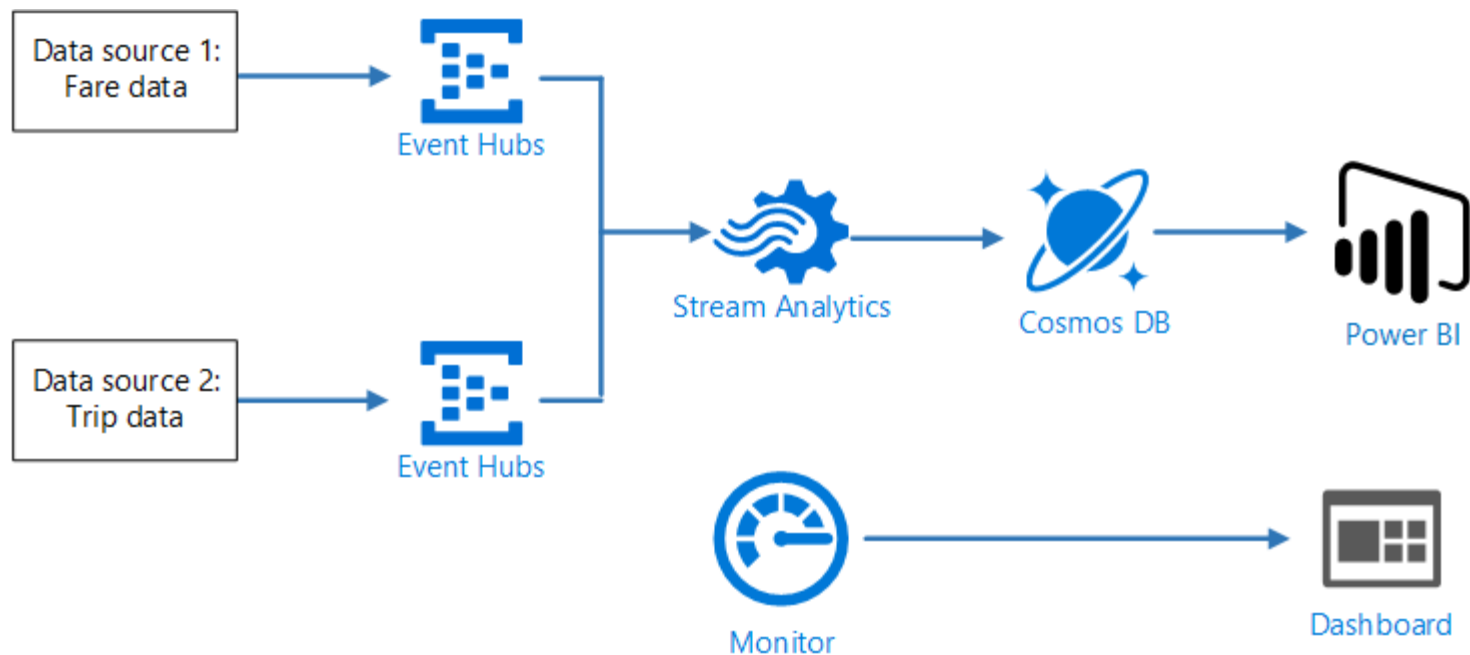


Design a Solution That Uses Azure Stream Analytics

- <https://docs.microsoft.com/en-us/azure/architecture/browse/#analytics>
 - <https://docs.microsoft.com/en-us/azure/architecture/reference-architectures/data/stream-processing-stream-analytics>




Azure Stream Analytics



<https://docs.microsoft.com/en-us/azure/architecture/reference-architectures/data/stream-processing-stream-analytics>



Azure Stream Analytics




<https://docs.microsoft.com/en-us/azure/architecture/example-scenario/data/fraud-detection>



Demo

- Azure Stream Analytics



Design for Data Security and Compliance

Design for Data Security and Compliance

- Design security for source data access
- Design security for data policies and standards



Design Security for Source Data Access

- Plan for secure endpoints (private/public)
- Choose the appropriate authentication mechanism, such as access keys, shared access, signatures (SAS), and Azure Active Directory (Azure AD)






Plan for Secure Endpoints

- Secure endpoints:
 - Azure Cosmos DB
 - Azure Storage Account
 - Azure Synapse Analytics
 - Azure Data Factory
 - Azure Databricks



Demo



- Securing endpoints
 - Azure Storage Account
 - Azure Cosmos DB
 - Azure Synapse Analytics



Design Security for Data Policies and Standards

- Design data encryption for data at rest and in transit
- Design for data auditing and data masking
- Design for data privacy and data classification
- Design a data retention policy
- Plan an archiving strategy
- Plan to purge data based on business requirements



Design Data Encryption for Data at Rest and in Transit

- Data encryption:
 - Azure Cosmos DB
 - Azure Storage Account
 - Azure Synapse Analytics





Azure SQL Database Security

- Transparent Data Encryption (TDE)
- Always Encrypted
- Dynamic Data Masking






Azure Policy

- Azure Policy helps to enforce organizational standards and to assess compliance at-scale.



Demo




- Azure Storage Account data lifecycle




Demo

- Azure Synapse data masking




Demo

- Azure Synapse data classification



Demo

- Azure Policy for data services



The Exam

Questions in DP-201

- Multiple choice
- Drag and drop
- Scenario based
- No hands-on labs (as of December 7th, 2020)



DP-201

- Exam DP-201 :

<https://docs.microsoft.com/en-us/learn/certifications/exams/dp-201>

- Skills measured :

<https://query.prod.cms.rt.microsoft.com/cms/api/am/binary/RE3VRMb>



Candidates for this exam are Microsoft Azure data engineers who collaborate with business stakeholders to identify and meet the data requirements to design data solutions that use Azure data services.

Azure data engineers are responsible for data-related design tasks that include designing Azure data storage solutions that use relational and non-relational data stores, batch and real-time data processing solutions, and data security and compliance solutions.

Candidates for this exam must design data solutions that use the following Azure services: Azure Cosmos DB, Azure Synapse Analytics, Azure Data Lake Storage, Azure Data Factory, Azure Stream Analytics, Azure Databricks, and Azure Blob storage.

Part of the requirements for: [Microsoft Certified: Azure Data Engineer Associate](#)

Related exams: [1 related exam](#)

Important: [See details](#)

[Go to Certification Dashboard](#)

Schedule exam

Exam DP-201: Designing an Azure Data Solution

Languages: English, Japanese, Chinese (Simplified), Korean

Retirement date: none

This exam measures your ability to accomplish the following technical tasks: design Azure data storage solutions; design data processing solutions; and design for data security and compliance.

[Schedule exam](#)

United States

\$165 USD*

Price based on the country in which the exam is proctored.

[Official practice test](#) for Designing an Azure Data Solution

All objectives of the exam are covered in depth so you'll be ready for any question on the exam.

My Profile

Exam Discounts

Verify exam discount eligibility

For Microsoft employees

Microsoft employees are eligible for discounted exams. The discount will be reflected at the end of the checkout process. For MOS exams at Certiport, please request a voucher through the Microsoft Employee Voucher Portal.

To verify you are a Microsoft employee, link your Microsoft work account (alias@microsoft.com).

Link account

For Microsoft event attendees

If you recently attended a Microsoft event, you may be eligible for a discounted Microsoft Certification exam. To check eligibility, select an event you attended and verify the account used to register for the event. [Terms and Conditions](#) apply.

Microsoft Ignite 2019, Orlando

Verify account

Continue scheduling exam

Proceed to the Pearson VUE website to complete the exam scheduling process.

Go to Pearson VUE

Contact us

Privacy & Cookies

Terms of use

Trademarks

Accommodations

© Microsoft 2020



Select exam options

DP-200: Implementing an Azure Data Solution

All fields are required.

How do you want to take your exam? [Exam delivery option descriptions](#)

- ☐ At a local test center
- ☒ At my home or office
- ☐ I have a Private Access Code

Are you going to be testing on this device and network?

If so, perform a quick pre-check to verify compatibility of your device and network before planning to take this exam in your home or office.

If you skip, be sure to do a full system test before test day to avoid lost exam fees and launch delays.

[Run pre-check](#)

[Next](#)





System check - Checking your requirements



Microphone

Default - Microphone (SI)



Internet speed



Webcam

Integrated Webcam (0c

Next



Course Repository

<https://github.com/zaalion/oreilly-dp-200-201>





Q&A



O'REILLY[®]

Thank you!

Reza Salehi

@zaalion

